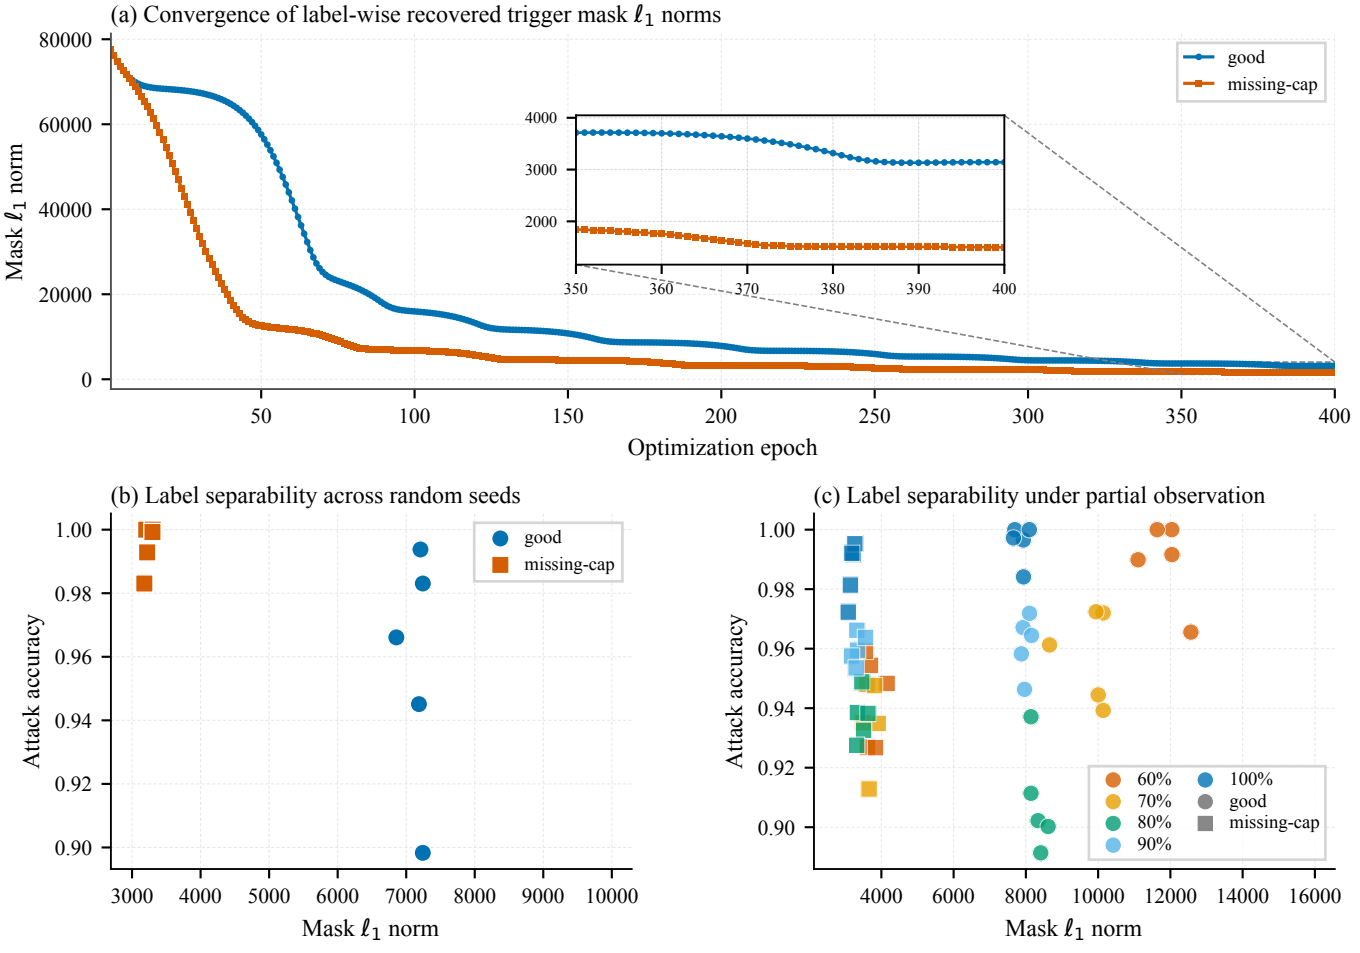
(a) Convergence of label-wise recovered trigger mask ℓ_1 norms
Mask ℓ_1 norm
Optimization epoch
good
missing-cap
(b) Label separability across random seeds
Attack accuracy
Mask ℓ_1 norm
good
missing-cap
(c) Label separability under partial observation
Attack accuracy
Mask ℓ_1 norm
60%
70%
80%
90%
100%
good
missing-cap